

[← Back to Blogs](#)

# AI Agent Best Practices: 12 Lessons from AI Pair Programming for Developers

June 1, 2025 • 7 min read

AI Coding    Pair Programming    Productivity    Software Engineering

AI Agent    Developer Best Practices    Workflow Optimization



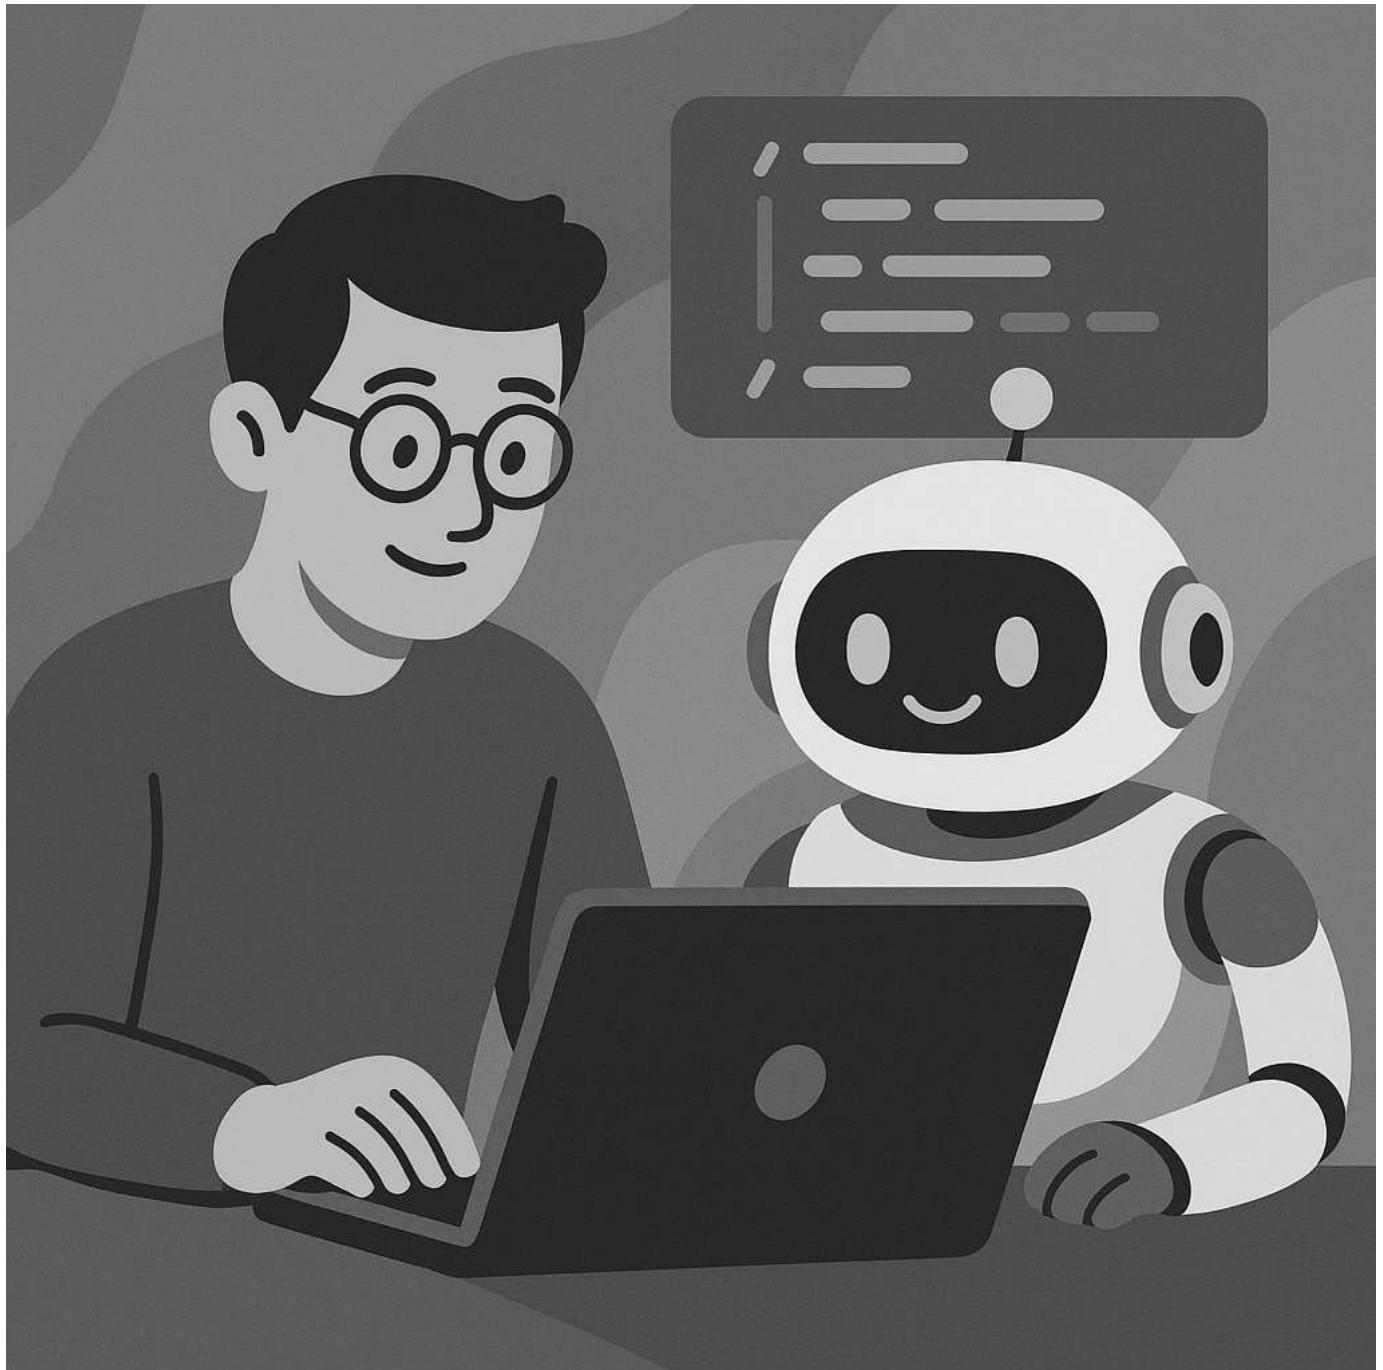
## We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)



After 6 months of daily AI pair programming across multiple codebases, here's what actually moves the needle. Skip the hype this is what works in practice.

## TL:DR



### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

## Prompt Engineering:

- Keep prompts short and specific context bloat kills accuracy
- Ask for step-by-step reasoning before code
- Use file references (@path/file.rs:42-88) not code dumps

## Context Management:

- Re-index your project after major changes to avoid hallucinations
- Use tools like gitingest.com for codebase summaries
- Use Context7 MCP to stay synced with latest documentation
- Treat AI output like junior dev PRs review everything

## What Doesn't Work:

- Dumping entire codebases into prompts
- Expecting AI to understand implicit requirements
- Trusting AI with security-critical code without review

# 1. Start With a Written Plan (Seriously, Do This First)

Ask your AI to draft a **Markdown plan** of the feature you're building. Then make it better:

### 1. Ask clarifying questions about edge cases

#### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

Write a plan for adding rate limiting to our API. Include:

- Which endpoints need protection
- Storage mechanism for rate data
- Error responses and status codes
- Integration points with existing middleware

Now critique this plan. What did you miss?

---

## 2. Master the Edit-Test Loop

This is TDD but with an AI doing the implementation:

1. **Ask AI to write a failing test that captures exactly what you want**
2. **Review the test yourself** - make sure it tests the right behavior
3. **Then tell the AI: "Make this test pass"**
4. **Let the AI iterate** - it can run tests and fix failures automatically

The key is reviewing the test before implementation. A bad test will lead to code that passes the wrong requirements.

## 3. Demand Step-by-Step Reasoning

Add this to your prompts:



### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

# 4. Stop Dumping Context, Start Curating It

Large projects break AI attention. Here's how to fix it:

## Use gitingest.com for Codebase Summaries

1. Go to gitingest.com
2. Enter your repo URL (or replace "github.com" with "gitingest.com" in any GitHub URL)
3. Download the generated text summary
4. Reference this instead of copy-pasting files

**Instead of:** Pasting 10 files into your prompt

**Do this:** "See attached codebase\_summary.txt for project structure"

## For Documentation: Use Context7 MCP or Alternatives for Live Docs

Context7 MCP keeps AI synced with the latest documentation by presenting the "Most Current Page" of your docs.

**When to use:** When your docs change frequently, reference the MCP connection rather than pasting outdated snippets each time.

### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

- Use meaningful commit messages: they help AI understand change context

## 6. Keep Prompts Laser-Focused

Bad: "Here's my entire codebase. Why doesn't authentication work?"

Good: "`asrc/auth.rs` line 85 panics on `None` when JWT is malformed. Fix this and add proper error handling."

Specific problems get specific solutions. Vague problems get hallucinations.

Use your code's terminology in prompts: reference the exact identifiers from your codebase, not generic business terms. For example, call `createOrder()` and `processRefund()` instead of 'place order' or 'issue refund', or use `UserEntity` rather than 'account'. This precision helps the AI apply the correct abstractions and avoids mismatches between your domain language and code.

## 7. Re-Index After Big Changes

If you're using AI tools with project indexing, rebuild the index after major refactors. Out-of-date indexes are why AI "can't find" functions that definitely exist.

Most tools auto-index, but force a refresh when things seem off.

### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

## Benefits:

- AI sees the current file state, not a stale snapshot
- Smaller token usage = better accuracy
- Less prompt clutter

**Note:** Syntax varies by tool (Forge uses `(@)`, some use `(#)`, etc.)

## 9. Let AI Write Tests, But You Write the Specs

Tell the AI exactly what to test:

For the new `validate\_email` function, write tests for:

- Valid email formats (basic cases)
- Invalid formats (no @, multiple @, empty string)
- Edge cases (very long domains, unicode characters)
- Return value format (should be `Result<()`, `ValidationError>`)

---

AI is good at generating test boilerplate once you specify the cases.

## 10. Debug with Diagnostic Reports



### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

#### 4. Propose 3 different debugging approaches

---

This forces the AI to think systematically instead of guess-and-check.

## 11. Set Clear Style Guidelines

Give your AI a brief system prompt:

Code style rules:

- Use explicit error handling, no unwraps in production code
- Include docstrings for public functions
- Prefer composition over inheritance
- Keep functions under 50 lines
- Use `pretty\_assertions` in test
- Be explicit about lifetimes in Rust
- Use `anyhow::Result` for error handling in services and repositories.
- Create domain errors using `thiserror`.
- Never implement `From` for converting domain errors, manually convert them

---

Consistent rules = consistent code quality.



### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

- Check for injection vulnerabilities
- Verify input validation
- Look for hardcoded secrets

### Performance Review:

- Watch for N+1 queries
- Check algorithm complexity
- Look for unnecessary allocations

### Correctness Review:

- Test edge cases manually
- Verify error handling
- Check for off-by-one errors

The AI is smart but not wise. Your experience matters.

## What Doesn't Work (Learn From My Mistakes)

### The "Magic Prompt" Fallacy

There's no perfect prompt that makes AI never make mistakes. Better workflows beat better prompts.

#### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)



[Accept All](#)

[Manage Settings](#)

[Deny](#)

AI is great at implementing your design but terrible at high-level system design.  
You architect, AI implements.

## Ignoring Domain-Specific Context

AI doesn't know your business logic, deployment constraints, or team conventions unless you tell it.

## Controversial Take: AI Pair Programming Is Better Than Human Pair Programming

For most implementation tasks.

AI doesn't get tired, doesn't have ego, doesn't argue about code style, and doesn't judge your googling habits. It's like having a junior developer with infinite patience and perfect memory.

But it also doesn't catch logic errors, doesn't understand business context, and doesn't push back on bad ideas. You still need humans for the hard stuff.

## Final Reality Check

AI coding tools can significantly boost productivity, but only if you use them



We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

The future of coding isn't human vs AI it's humans with AI vs humans without it.  
Choose your side wisely.

## Related Articles

- [Claude 4 Opus vs Grok 4: AI Model Comparison for Complex Coding Tasks]  
(/blog/slug: claude-4-opus-vs-grok-4-comparison-full)
- Claude Sonnet 4 vs Gemini 2.5 Pro Preview: AI Coding Assistant Comparison
- Forge Performance RCA: Root Cause Analysis of Quality Degradation on July 12, 2025
- MCP Security Prevention: Practical Strategies for AI Development - Part 2

Posted By



**Forge Team**



## Recent Blog Posts

### Forge v0.106.0 Release: Plan Progress Tracking and Reliability



#### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)

## Coding Agents Showdown: VSCode Forks vs. IDE Extensions vs. CLI Agents

The AI coding assistant landscape is fragmenting into three distinct ways to integrate AI into your...

Tushar



## Claude Sonnet 4 vs Kimi K2 vs Gemini 2.5 Pro: Which AI actually ships production code?

I ran Claude Sonnet 4, Kimi K2, and Gemini 2.5 Pro on the same Next.js app and measured cost,...

Amitesh Anand



## Graduating from Early Access: New Pricing Tiers Now Available

How our explosive early access growth shaped our pricing strategy and what's now available for...

Tushar



## Kimi K2 vs Grok 4: Which AI Model Codes Better?

A deep dive into Kimi K2 and Grok 4 for real-world coding, comparing their performance across bug...

Shrijal Acharya



## Kimi K2 vs Qwen-3 Coder: Testing Two AI Models on Coding Tasks

I tested Kimi K2 and Qwen-3 Coder on 13 Rust development tasks across a 38k-line codebase and 2...

Tushar



### We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)



## Grok 4 Initial Impressions: Is xAI's New LLM the Most Intelligent AI Model Yet?

A deep dive into Grok 4's benchmarks, architecture, and community impressions. Is xAI's latest LLM...

**Arindam Majumder**



## Claude 4 Opus vs Grok 4: Which Model Dominates Complex Coding Tasks?

I pitted Claude 4 Opus against Grok 4 in a series of challenging coding tasks. The results highlight...

**Tushar**



## Forge v0.98.0: Integrated Authentication and Developer Experience Improvements

Forge v0.98.0 release brings browser-based authentication, AI safety limits, and enhanced file...

**Forge Team**



Developers

Company



We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

[Accept All](#)

[Manage Settings](#)

[Deny](#)



## We Value Your Privacy

This website uses cookies to ensure you receive the best possible experience. [Learn More](#)

**Accept All**

**Manage Settings**

**Deny**